

## Case Study 2 – Customers and Transactions Data

The required tables CUSTOMERS and TRANSACTIONS are created using the To-do session of Hive basics session.

Creating the tables in hive.

```
hive> CREATE TABLE CUSTOMER(  
> custid INT,  
> fname STRING,  
> lname STRING,  
> age INT,  
> profession STRING)  
> row format delimited fields terminated by ',';  
OK  
Time taken: 20.704 seconds  
hive> LOAD DATA LOCAL INPATH "/home/acadgild/user_acadgild/custs.txt" into table CUSTOMER  
> ;  
Loading data to table default.customer  
OK  
Time taken: 5.528 seconds
```

```
hive> CREATE TABLE TRANSACTIONS(  
> txnno INT,  
> txndate STRING,  
> custno INT,  
> amount DOUBLE,  
> category STRING,  
> product STRING,  
> city STRING,  
> state STRING,  
> spendby STRING)  
> row format delimited fields terminated by ',';  
OK  
Time taken: 0.426 seconds  
hive> LOAD DATA LOCAL INPATH "/home/acadgild/user_acadgild/txn.txt" into table TRANSACTIONS;  
Loading data to table default.transactions  
OK  
Time taken: 1.995 seconds  
hive>
```

The data from both these tables can be seen below.

The data from both these tables can be seen below.

```
hive> select * from customer;  
OK  
101      Amitabh Bacchan 65      Actor  
102      Sharukh Khan 45      Doctor  
103      Akshay Kumar 38      Dentist  
104      Anubahv kumar 58      Business  
105      Pawan Trivedi 34      service  
106      Aamir Null 42      scientest  
107      Salman Khan 43      Surgen  
108      Ranbir Kapoor 26      Industrialist  
Time taken: 4.404 seconds, Fetched: 8 row(s)
```

```
hive> select * from transactions;  
OK  
97834 05/02/2018 101 965.0 Entertainment Movie Pune Maharashtra Daughter  
98396 12/01/2018 102 239.0 Food Grocery Patna Bihar Self  
34908 06/01/2018 101 875.0 Travel Air Bangalore Karnataka Spouse  
70958 17/02/2018 104 439.0 Food Restaurant Delhi Delhi Wife  
9874 21/01/2018 105 509.0 Entertainment Park Kolkata West Bengal NULL  
94585 19/01/2018 106 629.0 Rent House Hyderabad Telangana Self  
45509 20/01/2018 107 953.0 Travel Rail Chennai Tamil Nadu Brother  
7864 01/02/2018 108 569.0 Rent Parking Goa Goa Wife  
Time taken: 0.498 seconds, Fetched: 8 row(s)
```

## Objective 1:

Find out the number of transaction done by each customer.

Only the custno and the number of transactions can be queries only the transactions table.

```
select custno, count(*) from TRANSACTIONS group by custno;
```

```
hive> select custno, count(*) from TRANSACTIONS group by custno;  
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different  
tion engine (i.e. spark, tez) or using Hive 1.X releases.
```

```
MapReduce Jobs Launched:  
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU:  
Total MapReduce CPU Time Spent: 4 seconds 780 msec  
OK  
101 2  
102 1  
104 1  
105 1  
106 1  
107 1  
108 1  
Time taken: 59.681 seconds, Fetched: 7 row(s)
```

The custid, along with the customer's first name can be queried as follows.

```
select t.custno,c.fname, count(*) from CUSTOMER c, TRANSACTIONS t where  
c.custid=t.custno group by c.fname,t.custno;
```

```
MapReduce Jobs Launched:  
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU:  
Total MapReduce CPU Time Spent: 7 seconds 450 msec  
OK  
106 Aamir 1  
101 Amitabh 2  
104 Anubahv 1  
105 Pawan 1  
108 Ranbir 1  
107 Salman 1  
102 Sharukh 1  
Time taken: 136.514 seconds, Fetched: 7 row(s)
```

The custid, customer name and the count of times the customer has occurred in the Transactions table is shown, thus showing the number of transactions per each customer.

### Objective 2:

Create a new table called **TRANSACTIONS\_COUNT**. This table should have 3 fields - **custid**, **fname** and **count**.

```
create table TRANSACTIONS_COUNT(  
    custid INT,  
    fname STRING,  
    count INT)  
row format delimited fields terminated by ',';
```

```
hive> create table TRANSACTIONS_COUNT(  
    > custid INT,  
    > fname STRING,  
    > count INT)  
    > row format delimited fields terminated by ',';  
OK  
Time taken: 1.202 seconds
```

The table with the given column names is created.

### Objective 3:

Now write a hive query in such a way that the query populates the data obtained in Step 1 above and populate the table in step 2 above.

```
insert overwrite table TRANSACTIONS_COUNT select t.custno,c.fname, count(*)  
from CUSTOMER c, TRANSACTIONS t where c.custid=t.custno group by  
c.fname,t.custno;
```

```
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 7.77 sec  
Total MapReduce CPU Time Spent: 7 seconds 770 msec  
OK  
Time taken: 125.207 seconds  
hive> select * from TRANSACTIONS_COUNT;  
OK  
106      Aamir      1  
101      Amitabh   2  
104      Anubahv   1  
105      Pawan     1  
108      Ranbir    1  
107      Salman    1  
102      Sharukh   1  
Time taken: 0.54 seconds, Fetched: 7 row(s)
```

The data from step 1 is inserted into the newly created table **TRANSACTIONS\_COUNT**. The first line OK in the above screenshot is from when the command is run. And when queried the new table, the data is there.

#### Objective 4:

Now let's make the **TRANSACTIONS\_COUNT** table Hbase compliant. In the sense, use Ser Des and Storage handler features of hive to change the **TRANSACTIONS\_COUNT** table to be able to create a **TRANSACTIONS** table in Hbase.

For a table to be Hbase compliant and to load data into a hbase table from hive, a table has to be created in HBASE. Then, that table name can be specified in the serde properties command in Hive.

Creation of table in Hbase.

```
create 'TRANSACTIONS','txn_details'
```

```
hbase(main):001:0> create 'TRANSACTIONS','txn_details'
0 row(s) in 11.5330 seconds
=> Hbase::Table - TRANSACTIONS
```

The table name is **TRANSACTIONS** and column family is **txn\_details**. The columns from hive can be added into this column family.

Then, we have to create the Hive external table on top of HBase table that you want to populate.

```
CREATE EXTERNAL TABLE HBASE_TRANSACTIONS (custid INT, fname STRING, count INT)
  STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
  WITH SERDEPROPERTIES ("hbase.columns.mapping" =
":key,txn_details:fname,txn_details:count")
  TBLPROPERTIES("hbase.table.name"="transactions");
```

```
hive> CREATE EXTERNAL TABLE HBASE_TRANSACTIONS (custid INT, fname STRING, count INT)
> STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
> WITH SERDEPROPERTIES ("hbase.columns.mapping" = ":key,txn_details:fname,txn_details:count")
> TBLPROPERTIES("hbase.table.name"="TRANSACTIONS");
OK
Time taken: 5.076 seconds
hive> █
```

An external table is created. The **HBaseStorageHandler** is used in creation of the table because the table has to be Hbase compliant.

In the serde properties, the column mappings are specified as to which column in hive table is mapped to which column in which column family in Hbase.

Then in the table properties, the name of the Hbase table is specified.



### Objective 5:

Now insert the data in TRANSACTIONS\_COUNT table using the query in step 3 again, this should populate the Hbase TRANSACTIONS table automatically.

Populating the newly created external table.

```
insert into HBASE_TRANSACTIONS select t.custno,c.fname, count(*) from
CUSTOMER c, TRANSACTIONS t where c.custid=t.custno group by
c.fname,t.custno;
```

```
hive> insert into HBASE_TRANSACTIONS select t.custno,c.fname, count(*) from CUSTOMER c, TRANSACTIONS t where c.custid=t.custno
group by c.fname,t.custno;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execu
tion engine (i.e. spark, tez) or using Hive 1.X releases.
```

Once the command is successfully executed, the data can be seen in Hbase table too.

```
hbase(main):003:0> scan 'TRANSACTIONS'
ROW COLUMN+CELL
101 column=txn_details:count, timestamp=1531073229277, value=2
101 column=txn_details:fname, timestamp=1531073229277, value=Amitabh
102 column=txn_details:count, timestamp=1531073229277, value=1
102 column=txn_details:fname, timestamp=1531073229277, value=Sharukh
104 column=txn_details:count, timestamp=1531073229277, value=1
104 column=txn_details:fname, timestamp=1531073229277, value=Anubahv
105 column=txn_details:count, timestamp=1531073229277, value=1
105 column=txn_details:fname, timestamp=1531073229277, value=Pawan
106 column=txn_details:count, timestamp=1531073229277, value=1
106 column=txn_details:fname, timestamp=1531073229277, value=Aamir
107 column=txn_details:count, timestamp=1531073229277, value=1
107 column=txn_details:fname, timestamp=1531073229277, value=Salman
108 column=txn_details:count, timestamp=1531073229277, value=1
108 column=txn_details:fname, timestamp=1531073229277, value=Ranbir
7 row(s) in 0.6750 seconds
```

The table contents in hbase are displayed and the 7 rows can be seen in hbase too.

So, the table in Hbase is automatically populated when the hive table is populated.

#### Objective 6:

Now from the Hbase level, write the Hbase java API code to access and scan the TRANSACTIONS table data from java level.

```
package com.acadgild.cs2;

import java.io.IOException;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.hbase.HBaseConfiguration;
import org.apache.hadoop.hbase.util.Bytes;
import org.apache.hadoop.hbase.client.HTable;
import org.apache.hadoop.hbase.client.Result;
import org.apache.hadoop.hbase.client.ResultScanner;
import org.apache.hadoop.hbase.client.Scan;

public class ScanTable {
    public static void main(String args[]) throws IOException {
        Configuration config = HBaseConfiguration.create();
        @SuppressWarnings({ "deprecation", "resource" })
        HTable table = new HTable(config, "TRANSACTIONS");
        Scan scan = new Scan();
        scan.addColumn(Bytes.toBytes("txn_details"), Bytes.toBytes("count"));
        scan.addColumn(Bytes.toBytes("txn_details"), Bytes.toBytes("fname"));
        ResultScanner scanner = table.getScanner(scan);
        for (Result result = scanner.next(); result != null; result = scanner.next()) {
            String Row = Bytes.toString(result.getRow());
            String name = Bytes.toString(result.getValue("txn_details".getBytes(),
                "fname".getBytes()));
            String count = Bytes.toString(result.getValue("txn_details".getBytes(),
                "count".getBytes()));
            System.out.println(Row + "," + name + "," + count);
            scanner.close();
        }
    }
}
```

```
101,Amitabh,2
102,Sharukh,1
104,Anubahv,1
105,Pawan,1
106,Aamir,1
107,Salman,1
108,Ranbir,1
```

